

Muhammad Abdullah

Date of birth: 13/11/2005 | **Nationality:** Pakistani | **Phone number:** (+92) 3094157102 (Mobile) | **Email address:** abdullahse2211@gmail.com | **LinkedIn:** [Muhammadabdullahse2211](#) | **GitHub:** [abdullahse2211](#) | **My Portfolio :** [Muhammad Abdullah](#) | **Whatsapp Messenger:** +92 3094157102 | **Address:** Pakistan (Home)

EDUCATION AND TRAINING

02/11/2022 – 10/08/2026 Lahore, Pakistan
BACHELOR OF SOFTWARE ENGINEERING Lahore Garrison University, DHA, Lahore

Principal subjects: Machine Learning; Big Data Analytics; Data Structure and Algorithms; Linear Algebra; Calculus and Analytical Geometry; Probability and Statistics; Stochastic Processes, Operational Research, Discrete Structures, Theory of Automata, Programming Fundamentals, Object-Oriented Programming, Database Systems, Software Design and Architecture, Software Construction and Development, and Final Year Project in Medical Imaging Computer Vision. AI

Level in EQF EQF level 6

WORK EXPERIENCE

AI/ML ENGINEER (PART-TIME, 30 HRS/WEEK) – ALLYTICA – 01/09/2025 – 03/09/2026 – LAHORE, PAKISTAN

- Agentic GraphRAG & Knowledge Architectures:** Co-architected enterprise Agentic GraphRAG pipelines within a specialized 2-person AI engineering unit, integrating Neo4j Knowledge Graphs with LlamaIndex agents to reduce multi-hop retrieval hallucinations by 42%.
- Sovereign LLM Quantization & Edge Optimization:** Collaborated directly with a senior systems engineer to quantize open-weights LLMs/VLMs via 4-bit AWQ and deploy containerized vLLM microservices, reducing VRAM footprint by 60% and first-token latency by 55%.
- Multimodal Vision Engine Acceleration:** Spearheaded the fine-tuning of multimodal Vision-Language Models (Qwen2-VL) using LoRA under senior technical oversight, converting vision backbones into TensorRT FP16 execution engines to achieve 62 FPS.
- Enterprise Risk Analytics & Anomaly Detection:** Developed end-to-end statistical risk scoring and anomaly detection pipelines in PyTorch, maintaining a 0.91+ AUC on out-of-distribution enterprise transaction streams.
- Sovereign Cloud & Microservice Architecture:** Engineered secure REST API microservices using FastAPI and Docker in a pair-programming environment, implementing strict data residency guardrails for sovereign cloud compliance.

ACADEMIC & AI LAB RESEARCH PROJECTS

01/01/2026 – 22/05/2026
RetinaXplain - Hybrid CNN-Vision Transformer Framework for Diabetic Retinopathy Classification

- Final Year Project-Sole AI Researcher and Developer
- Data Pipeline & Preprocessing:** Processed 3,662 APTOS images across 5 DR classes using circular cropping and Ben Graham normalization, expanding data to ~16,000 augmented instances.
 - Model Training & Algorithmic Features:** Trained and evaluated Swin-T (85.27% accuracy, 0.7216 macro-F1), EfficientNet-B3, and ViT-B/16, engineering weighted soft-voting ensembles, clinical label remap,& agreement analysis.
 - App Deployment & Documentation:** Built a multi-page web app featuring real-time DR diagnostic reports, probability visualizers, session analytics, and downloadable JSON records; authored a 72-page technical report.

Links <https://github.com/abdullahse2211/diabetic-retinopathy-classification> | <https://youtu.be/MnS1vS5ZLqU?si=td3M1nW6vMmZRVf5>

07/01/2026 – 08/03/2026
BioNAFNet - Low-Light Image Enhancement | Student Deep Learning Researcher

- Novel Architecture & Loss Formulation: Designed a bio-inspired enhancement network trained from scratch with a composite Charbonnier, SSIM, color, and Weber-Fechner perceptual loss, surpassing 9 baselines with 28.67 dB PSNR on LOLv2 Real.
- Knowledge Distillation & Acceleration: Compressed model parameters from 154M to 24M via knowledge distillation (6.4x compression, 3.9x speedup) with zero perceptual degradation, enabling the student model to surpass the teacher by +0.29 dB.

Clinical Feature Engineering: Engineered 15 diagnostic features combining 9 CBC/HPLC parameters with 6 clinical indices (Mentzer, Shine & Lal, Green & King), elevating Silent Carrier diagnostic F1 to 0.92–0.955.

Robust Multi-Class ML Benchmarking: Evaluated 6 ML architectures (hybrid CNN-Transformer, 1D-CNN, SVM, RF, LR) for 5-class differential diagnosis, achieving 90.60% balanced accuracy and 0.906 weighted F1 under injected clinical noise and 10% feature dropout.

● SKILLS

ML and DL Frameworks: PyTorch, scikit-Learn, Hugging Face Transformers, TensorFlow | Gen AI: Diffusion Models (DDPM, DDIM), RAG, Bayesian Optimization (BoTorch) | Libraries and Tools: NumPy, Pandas, OpenCV, Matplotlib, Streamlit, Git, GitHub, Jupyter, Overleaf | AI Systems & MLOps: Qwen2-VL, LlamaIndex, GraphRAG, vLLM, TensorRT, 4-bit AWQ, Neo4j, LoRA. | Domains: Computer Vision, Image Processing, Time Series Forecasting, Generative AI, Data Analysis | Programming Languages: Python, MERN STACK, JavaScript (ES6+), SQL (MySQL), HTML5, CSS3

● CERTIFICATIONS

Certiport, a Pearson VUE Business, 13/01/2025
Information Technology Specialist: JavaScript

International industry certification in JavaScript. Credential ID: mTp4-DwVM

Mode of learning: Presential

Link <https://verify.certiport.com/>

National Vocational & Technical Training Commission (NAV TTC), 16/12/2024
JavaScript Full-Stack Development (MERN)

Completed full-stack JavaScript development training covering the MERN technology stack: MongoDB, Express.js, React, and Node.js. Grade A+.

Mode of learning: Project based

University of Engineering and Technology, Lahore National Centre of Artificial Intelligence, 08/12/2023
Prompt Engineering - Certificate of Completion

Completed practical training in prompt design and refinement for generative AI

Mode of learning: Presential

IBM; Coursera, 18/11/2023
IBM Web Development

Completed IBM training in foundational web-development concepts, tools, and practices.

Mode of learning: Online

Link <https://coursera.org/verify/E8X42H6HASJL>

● LANGUAGE SKILLS

Mother tongue(s): **URDU**

Other language(s): **ENGLISH**